

# Segment Routing


Network Infrastructures A.A. 2020/21

- Introduction to Segment Routing
- Segment Routing Policy
- SRv6 and Segment Routing extension Header
- Network Programming


- Introduction to Segment Routing
- Segment Routing Policy
- SRv6 and Segment Routing extension Header
- Network Programming

- Segment Routing (SR) leverages the **source routing paradigm**
- A node steers a packet through an ordered list of instructions, called **segments**
- A segment can represent any instruction, topological or service-based
- A segment can have a semantic:
  - local to an SR node
  - global within an SR domain
- SR allows to enforce a flow through any topological path while maintaining **per-flow state only at the ingress nodes to the SR domain**


# SR terminology




# SR idea



# SR idea




# SR idea





# SR idea



- The SR architecture supports any type of control-plane: distributed, centralized or hybrid.
- In a **distributed scenario**:
  - the segments are allocated and signaled by OSPF or BGP
  - a node individually decides to steer packets on a source-routed policy
  - a node individually computes the source-routed policy
- In a **centralized scenario**:
  - the segments are allocated and instantiated by an SR controller
  - the SR controller decides which nodes need to steer which packets on which source-routed policies
  - the SR controller computes the source-routed policies
- A **hybrid scenario** complements a base distributed control-plane with a centralized controller
  - for example, when the destination is outside the IGP domain, the SR controller may compute a source-routed policy on behalf of an IGP node


# SR data plane

---

- The SR architecture can be instantiated on various dataplanes
  - SR over MPLS (SR-MPLS)
  - SR over IPv6 (SRv6)
- Segment Routing can be directly applied to the MPLS architecture with no change on the forwarding plane
- Segment Routing can be applied to the IPv6 architecture with a new type of routing header called the SR header (SRH)

# SR forwarding


- SR nodes have a forwarding table that specifies the operation to perform on a received packet



- Three different operations:
  - PUSH**: the instruction consisting of the insertion of a segment at the top of the segment list

# SR forwarding


- SR nodes have a forwarding table that specifies the operation to perform on a received packet



- Three different operations:
  - **PUSH**: the instruction consisting of the insertion of a segment at the top of the segment list
  - **NEXT**: when the active segment is completed, NEXT is the instruction consisting of the inspection of the next segment

# SR forwarding

- SR nodes have a forwarding table that specifies the operation to perform on a received packet



- Three different operations:
  - PUSH**: the instruction consisting of the insertion of a segment at the top of the segment list
  - NEXT**: when the active segment is completed, NEXT is the instruction consisting of the inspection of the next segment
  - CONTINUE**: the active segment is not completed and hence remains active

# Global and Local segments

---

- **SR Global Block (SRGB)**: the set of global segments in the SR Domain
- **SR Local Block (SRLB)**: local property of an SR node
- Global Segment
  - the instruction associated to the segment is defined at the SR Domain level
  - a topological shortest-path segment to a given destination within an SR domain is a typical example of a global segment
- Local Segment
  - the instruction associated to the segment is defined at the node level


# Type of segments

---


- Segments are generally advertised in the network by means of an IGP protocol
- **IGP-Prefix segment**
  - forward the packet along the path computed using the routing algorithm
- **IGP-Node segment**
  - is an IGP-Prefix Segment which identifies a specific router
- **IGP-Adjacency segment**
  - forward the packet over a unidirectional adjacency (local segment)




# Type of segments: some examples




# Type of segments: some examples




# Type of segments: some examples



# Type of segments: some examples



# Type of segments: some examples



- Introduction to Segment Routing
- Segment Routing Policy
- SRv6 and Segment Routing extension Header
- Network Programming

- An **SR Policy** is identified through the tuple

**<headend, color, endpoint>**

- The **headend** is the node where the policy is instantiated/implemented
- The **endpoint** indicates the destination of the policy
  - headend and endpoint are specified as an IPv4 or IPv6 address
- The **color** is a 32-bit numerical value that associates the SR Policy with an intent (e.g., low-latency)

# Candidate Path and Segment List

---

- An SR Policy is associated with one (or more) candidate path
- A **candidate path** is itself **associated with** a **Segment-List** (SID-List)
  - a SID-List represents a specific source-routed way to send traffic from the head-end to the endpoint of the corresponding SR policy
- A candidate path is either **dynamic** or **explicit**
- A headend may be informed about a candidate path for an SR Policy by various means including:
  - local configuration
  - PCE



# SR Policy: summary

---

SR policy POL1 <headend, color, endpoint>

## Candidate-path CP1

*Preference 200*

Weight W1, SID-List1 <SID11...SID1i>

Weight W2, SID-List2 <SID21...SID2j>

## Candidate-path CP2

*Preference 100*

Weight W3, SID-List3 <SID31...SID3i>

Weight W4, SID-List4 <SID41...SID4j>

- An SR headend maintains the **Segment Routing Traffic Engineering Database (SRTE-DB)**
- The SRTE-DB is used to validate explicit candidate paths and compute dynamic candidate paths
- It includes the following information:
  - Regular IGP information (topology, IGP metrics)
  - Extended TE Link attributes (such as latency, loss, TE metric)
  - Inter-Domain Topology information
  - Segment Routing information (such as SRGB, Prefix-SIDs, Adj-SIDs, Peering SID SRv6 SID)


# Dynamic Candidate Path

---

- A dynamic candidate path is specified as an **optimization objective** and **constraints**
  - eg. minimize delay avoiding link l
- The headend of the policy leverages its SRTE-DB to compute a SID-List that fits this optimization problem
  - re-computes any time the inputs to the problem change
- When local computation is not possible, the head-end may send path computation request to a PCE


# Binding SID

- An SR Policy installs a **BSID** entry in the forwarding plane with the action of **steering the packets matching this entry to the selected path**
- Let us assume that headend **H** has a valid SR Policy **P** of SID-List **<S1, S2, S3>** and BSID **B**




- **H** has steered the packet in the SR policy **P**


# SRTE Policy: an example



# SRTE Policy: an example




# SRTE Policy: an example



# SRTE Policy: an example

option 1






# SRTE Policy: an example

option 2


install BSID:  $\langle D, D2E, DCI3 \rangle$   
advertise BSID in data center 1 as “low latency policy”




- Introduction to Segment Routing
- Segment Routing Policy
- SRv6 and Segment Routing extension Header
- Network Programming

# SRH instantiation

- The source based routing model in case of SR over IPv6 dataplane (SRv6) is realized through the instantiation of the **Segment Routing Header (SRH)**
- The SRH is added to the packet by its source:
  - At the **node originating the packet** (host, server)
  - At the **ingress node** of an SR domain




# Examples of SRv6: Service Provider Network



- an IPv6 packet received at ingress, is:
  - **classified** according to network operator policies
  - **encapsulated** with an outer header with an SRH applied to the incoming packet

# Examples of SRv6: Overlay Network



- The SRH originated by the overlay can only contain address/segment under the administration of the overlay

# Segment Routing Extension Header (SRH)

- A new type of the Routing Header is defined
  - the Segment Routing Header (SRH)

next header	hdr ext len	routing type	segments left
last entry	flags	tag	
segments list [0] (128 bit IPv6 address)			
.....			
segments list [n] (128 bit IPv6 address)			
optional type length value objects (variable)			

# Segment Routing Extension Header (SRH)

- **Next Header:** Identifies the type of header immediately following the SRH
- **Hdr Ext Len:** length of the SRH header in 8-octet units
- **Routing Type:** TBD, to be assigned by IANA (suggested value: 4)
- **Segments Left:** it contains the index, in the Segment List, of the next segment to inspect
  - Segments Left is decremented at each segment
- **Last Entry:** contains the index, in the Segment List, of the last element of the Segment List
- **Flags:** 8 bits of flags
- **Tag:** tag a packet as part of a class or group of packets (packets sharing the same set of properties)
- **Segment List[n]:** 128 bit IPv6 addresses representing the nth segment in the Segment List
- **Type Length Value (TLV)**

- Type Length Value (TLV) contain optional information that may be used by the node identified in the DA of the packet
  - Multiple TLVs may be encoded in the same SRH
- The following TLVs are defined:
  - **Ingress Node TLV**: identifies the node this packet traversed when entered the SR domain
  - **Egress Node TLV**: identifies the node this packet is expected to traverse when exiting the SR domain
  - **NSH Carrier TLV**



# SRH processing

- For the SRH holds the following property:
  - Only the router whose address is in the DA field of the packet header **MUST** inspect the SRH
- Segment Routing in IPv6 networks implies that the segment identifier is moved into the DA of the packet



- The DA of the packet changes at each segment termination/completion

# My Local SID Table

- An SRv6-capable node N maintains a **MyLocalSID Table**
- This table contains all the **local SRv6 segments** explicitly instantiated at node N
- N is the parent node for these SID's
- Every SRv6 local SID instantiated has a specific instruction bounded to it

MyLocalSID Table


SID	instruction
...	...

- Segment Routing architecture defines a **segment as an instruction** or, more generally, a set of instructions (function)
- Two SRv6 basic functions are:
  - **End**
    - the endpoint (End) function is the base of the source routing paradigm
    - it consists of **updating the DA with the next segment** and forward the packet accordingly
  - **End.X**
    - the endpoint **layer-3 cross-connect** function

# Endpoint Function (End)

- When a node receives a packet destined to DA "S" and:
  - S is an entry in the MyLocalSID table
  - the function associated with S is "End"

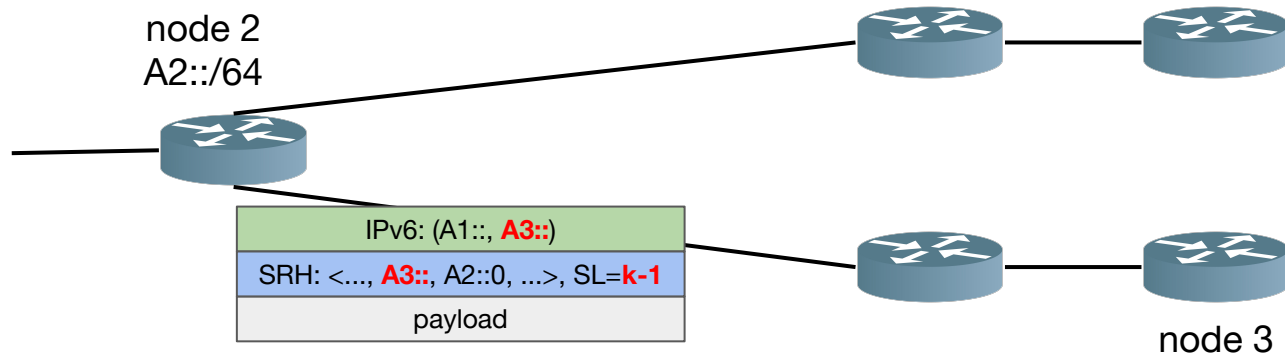
1. **IF** SegmentsLeft > 0 **THEN**
2.   decrement SL
3.   update the IPv6 DA with SRH[SL]
4.   FIB lookup on updated DA
5.   forward accordingly to the matched entry
6. **ELSE**
7.   drop the SRH



# Endpoint Function (End)

- When a node receives a packet destined to DA "S" and:
  - S is an entry in the MyLocalSID table
  - the function associated with S is "End"


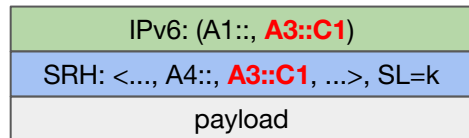
1. **IF** SegmentsLeft > 0 **THEN**
2.     decrement SL
3.     update the IPv6 DA with SRH[SL]
4.     FIB lookup on updated DA
5.     forward accordingly to the matched entry
6. **ELSE**
7.     drop the SRH



# Endpoint with Layer-3 cross-connect (End.X)

- When a node receives a packet destined to DA "S" and:
  - S is an entry in the MyLocalSID table
  - the function associated with S is "End.X"

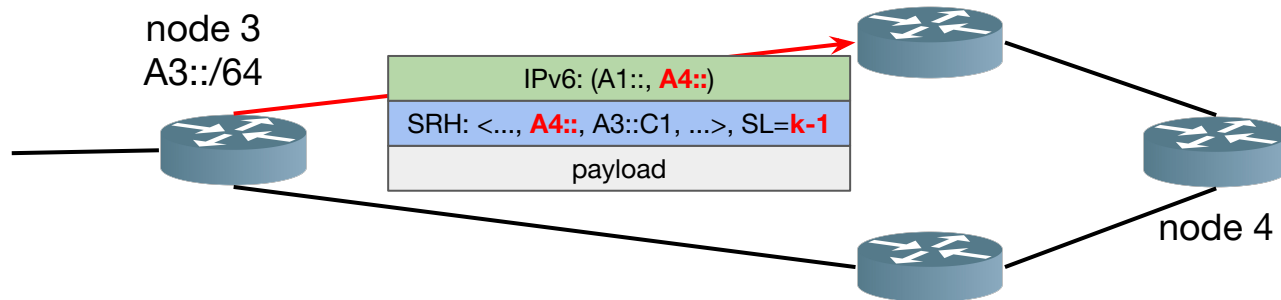
1. **IF** SegmentsLeft > 0 **THEN**
2.   decrement SL
3.   update the IPv6 DA with SRH[SL]
4.   forward to layer-3 adjacency bound to the SID "S"
5. **ELSE**
6.   drop the SRH



# Endpoint with Layer-3 cross-connect (End.X)

- When a node receives a packet destined to DA "S" and:
  - S is an entry in the MyLocalSID table
  - the function associated with S is "End.X"

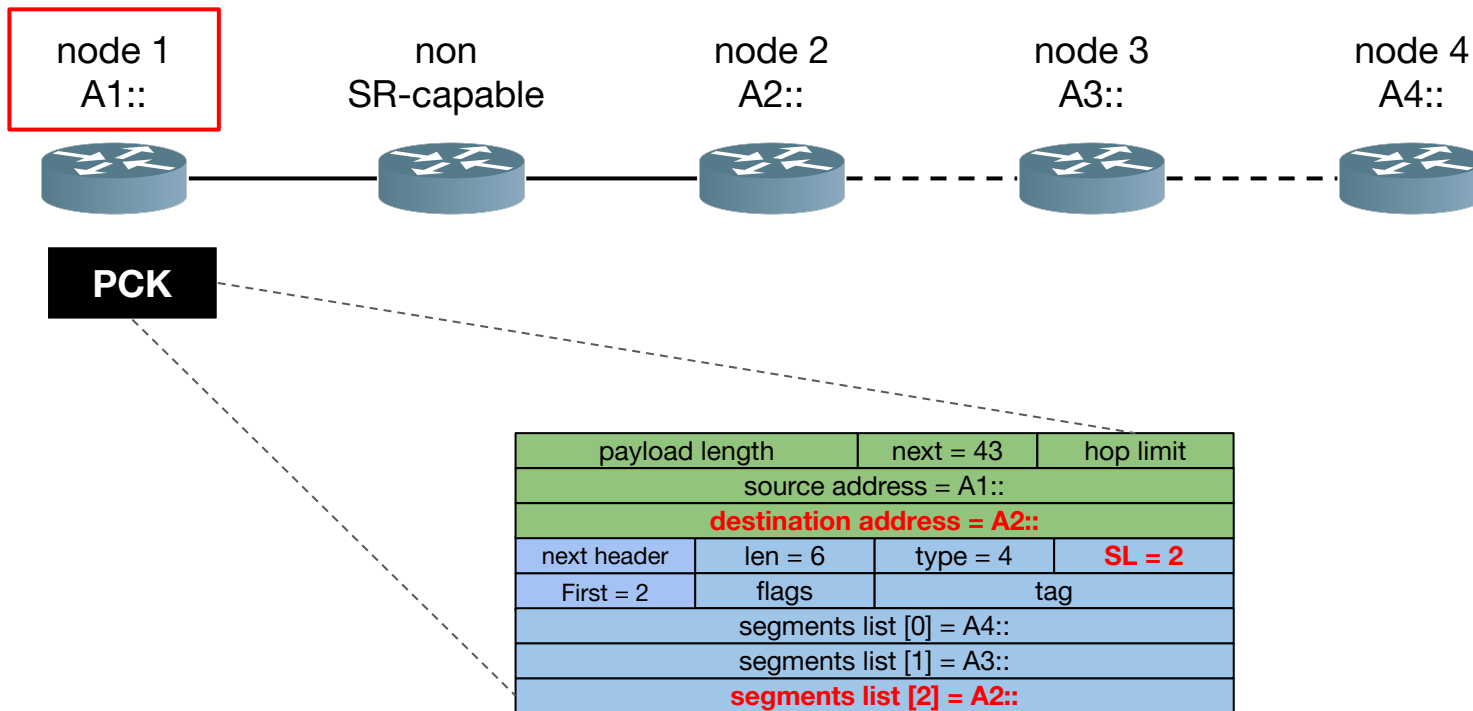
1. **IF** SegmentsLeft > 0 **THEN**
2.     decrement SL
3.     update the IPv6 DA with SRH[SL]
4.     forward to layer-3 adjacency bound to the SID "S"
5. **ELSE**
6.     drop the SRH




- Source Node is SR-capable
- SRH is created with:
  - Segment List in reversed order of the path
    - Segment List [0] is the **LAST** segment
    - Segment List [n-1] is the **FIRST** segment
    - Segment left is set to **n-1**
    - First segment is set to **n-1**
- IP DA is set to the first segment
- Packet is sent according to the IP DA
  - Normal IPv6 forwarding



# Source Node



# Non-SR Transit Node




- Plain IPv6 forwarding
- Solely based on IPv6 destination address
- No SRH inspection or update

# SR Segment Endpoints

---

- **SR Endpoints:** SR-capable nodes whose address is in the IP DA
- SR Endpoints inspect the SRH and do:
  - **IF Segments Left > 0, THEN**
    - Decrement Segments Left (-1)
    - Update DA with Segment List [Segments Left]
    - Forward according to the new IP DA


# SR Segment Endpoints



PCK

vers	traffic class	flow label	
payload length		next = 43	hop limit
source address = A1::			
destination address = A3::			
next header	len = 6	type = 4	SL = 1
first = 2	flags	tag	
segments list [0] = A4::			
segments list [1] = A3::			
segments list [2] = A2::			
Payload			

# SR Segment Endpoints



PCK

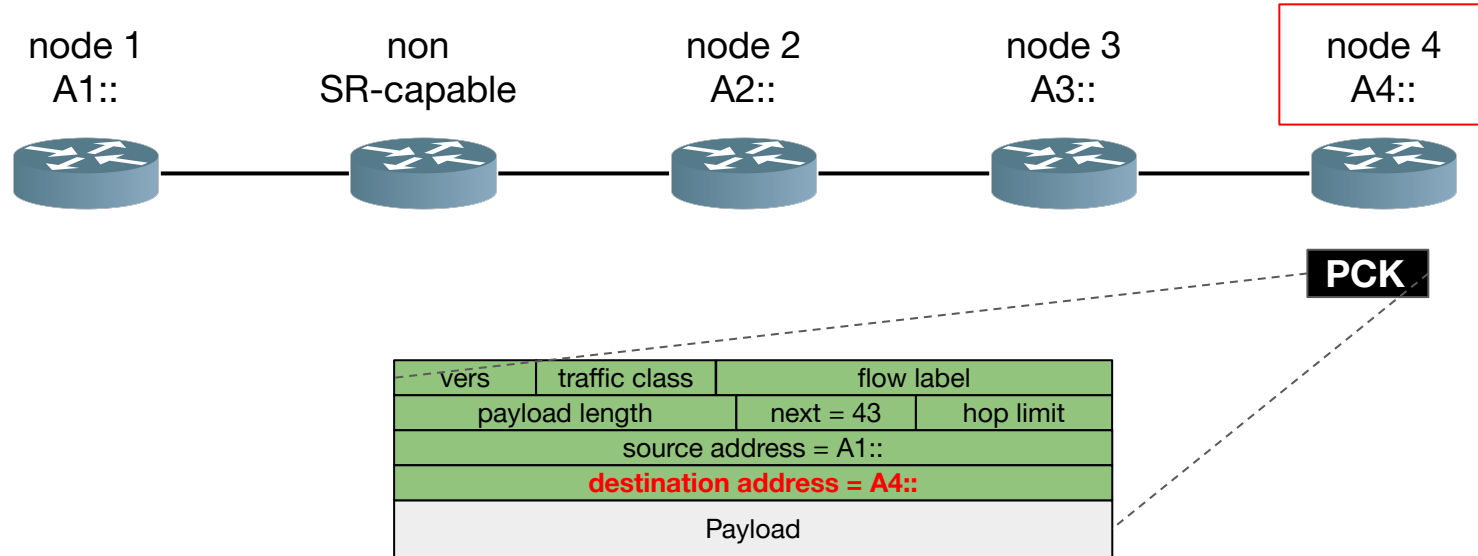
vers	traffic class	flow label	
payload length		next = 43	hop limit
source address = A1::			
destination address = A4::			
next header	len = 6	type = 4	SL = 0
first = 2	flags	tag	
segments list [0] = A4::			
segments list [1] = A3::			
segments list [2] = A2::			
Payload			

# SR Segment Endpoints

---

- SR Endpoints: SR-capable nodes whose address is in the IP DA
- SR Endpoints inspect the SRH and do:
  - IF Segments Left  $> 0$ , THEN
    - Decrement Segments Left (-1)
    - Update DA with Segment List [Segments Left]
    - Forward according to the new IP DA
  - ELSE (Segments Left = 0)
    - Remove the IP and SR header
    - Process the payload

# SR Segment Endpoints



- Introduction to Segment Routing
- Segment Routing Policy
- SRv6 and Segment Routing extension Header
- Network Programming



# SRv6 Segment Format

<i>locator</i>	<i>function</i>	<i>argument</i>
1111:2222:3333:4444	5555:6666	7777:8888



- An SRv6 local SID is logically represented as LOC:FUNCT
  - LOC is the L most significant bits
  - FUNCT is the 128-L least significant bits
- L is called the **locator** length and is flexible
- Most often the **LOC part of the SID is routable** and leads to the node which owns that SID
- The **FUNCT** part of the SID is an opaque identification of a **local function** bound to the SID
- A function may require additional arguments
  - the SRv6 Local SID will have the form LOC:FUNCT:ARGS

- Each entry of the MyLocalSID Table indicates the function associated with the local SID
- In practice, any function can be attached to a local SID
  - a node N can bind a SID to a local VM which can apply any complex function on the packet
- Some examples:
  - End
  - End.X
  - End.B6

# End.B6

- Endpoint bound to an SRv6 Policy
- When N receives a packet destined to S and S is a local End.B6 SID, N does:

1. **IF** NH=SRH **and** SL > 0
2. do not decrement SL nor update the IPv6 DA with SRH[SL]
3. insert a new SRH
4. set the IPv6 DA to the first segment of the SRv6 Policy
5. forward according to the first segment of the SRv6 Policy
6. **ELSE**
7. drop the SRH




MyLocalSID Table

SID	instruction
...	...
A4::20	End.X <S1, S2, S3>
...	...




# SID allocation for illustration purpose


- Node **K** advertises prefix **AK::/64**
- The **function** is encoded in the **last 64 bits**
  - **0** denotes the **End** function
  - **CJ** denotes the **End.X** function on **link CJ**




# A1::0 and then A5::0




# A1::0 and then A5::57



- 50 msec protection upon local link, node or SRLG failure

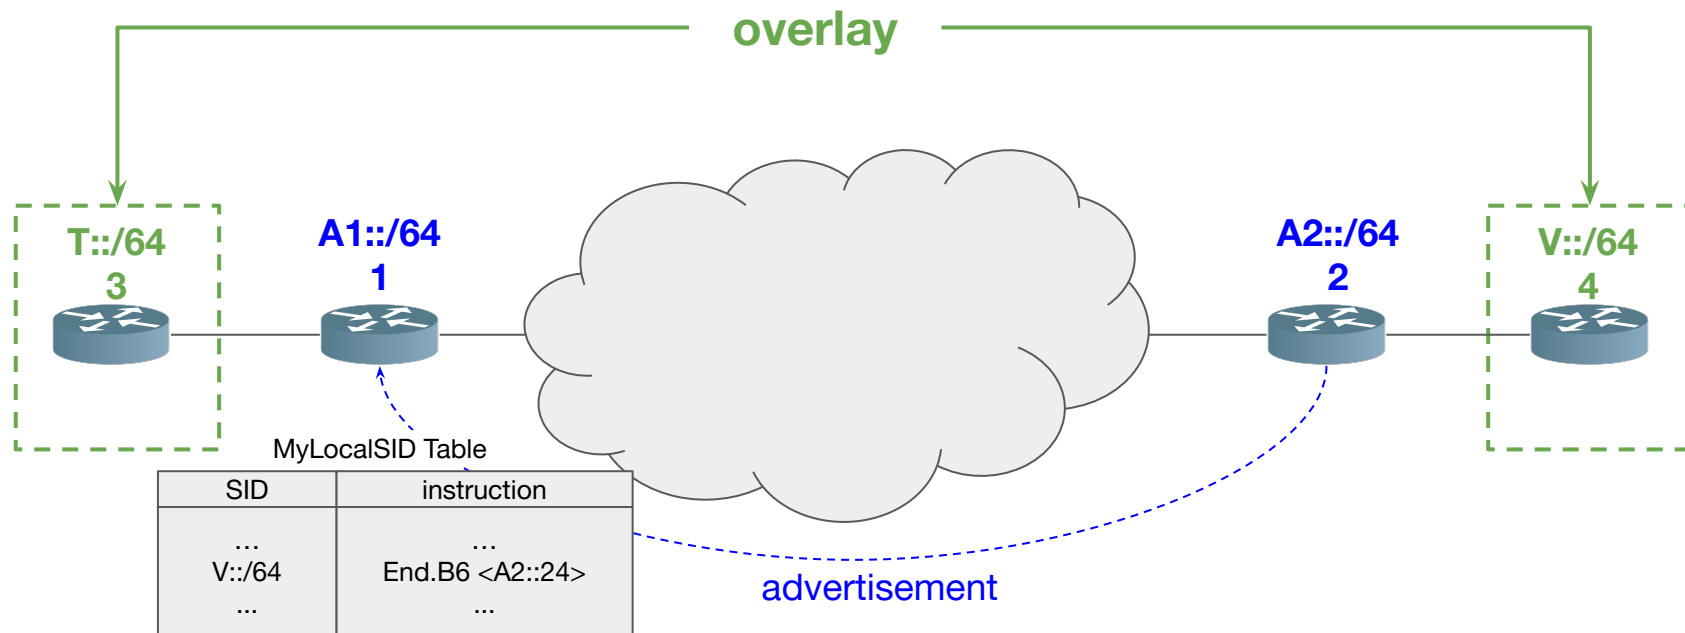


- 50 msec protection upon local link, node or SRLG failure






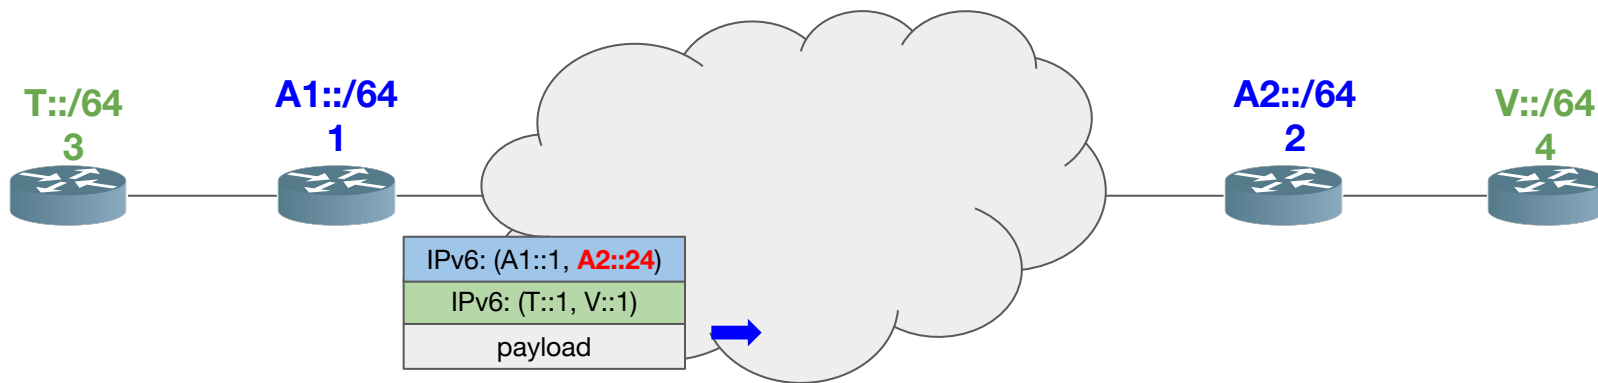
# Overlay




# Overlay



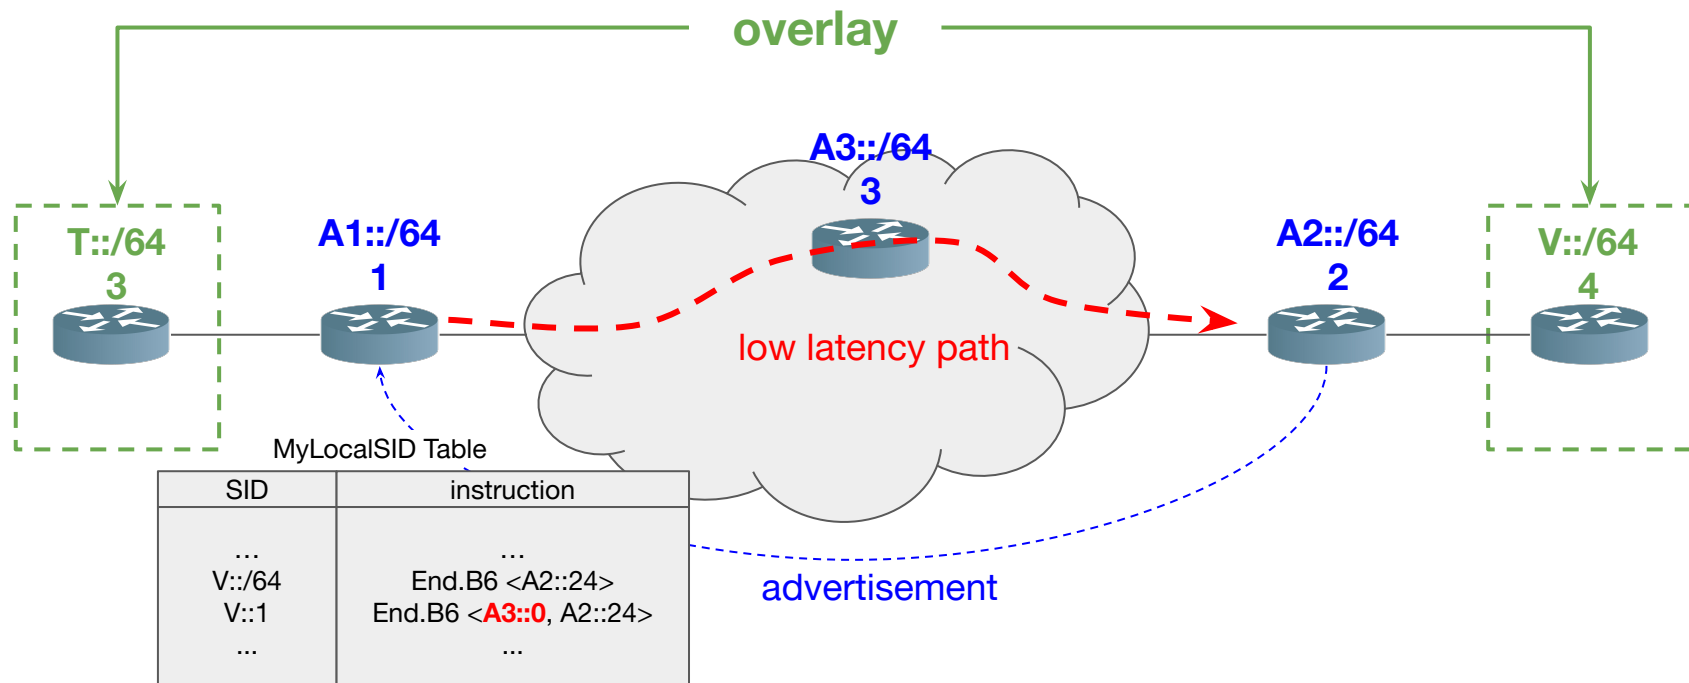
# Overlay




# Overlay



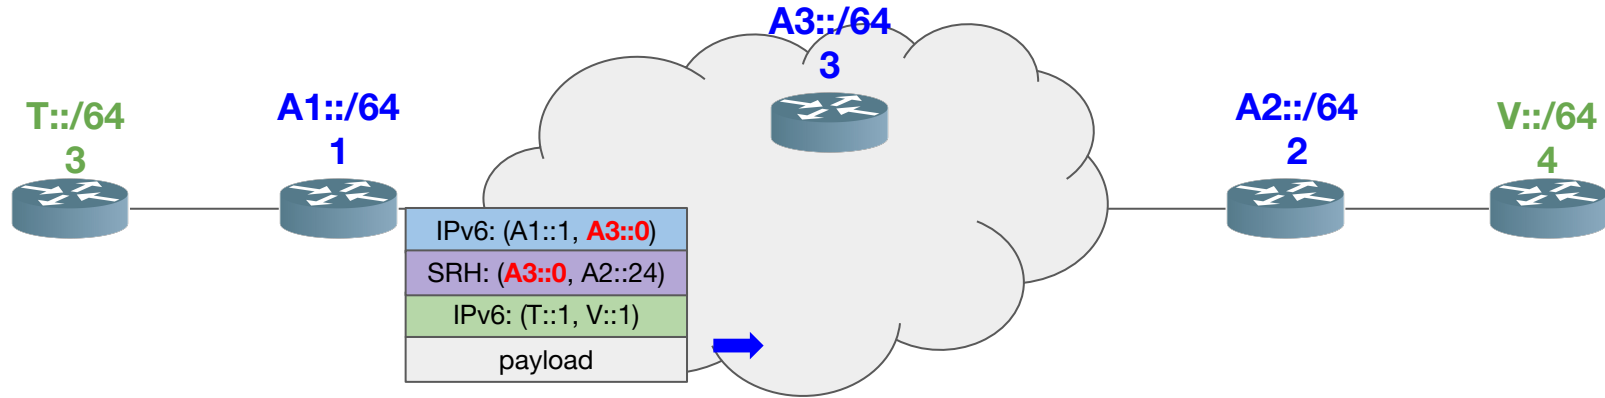
# Overlay with underlay SLA




# Overlay with underlay SLA



# Overlay with underlay SLA




# Overlay with underlay SLA

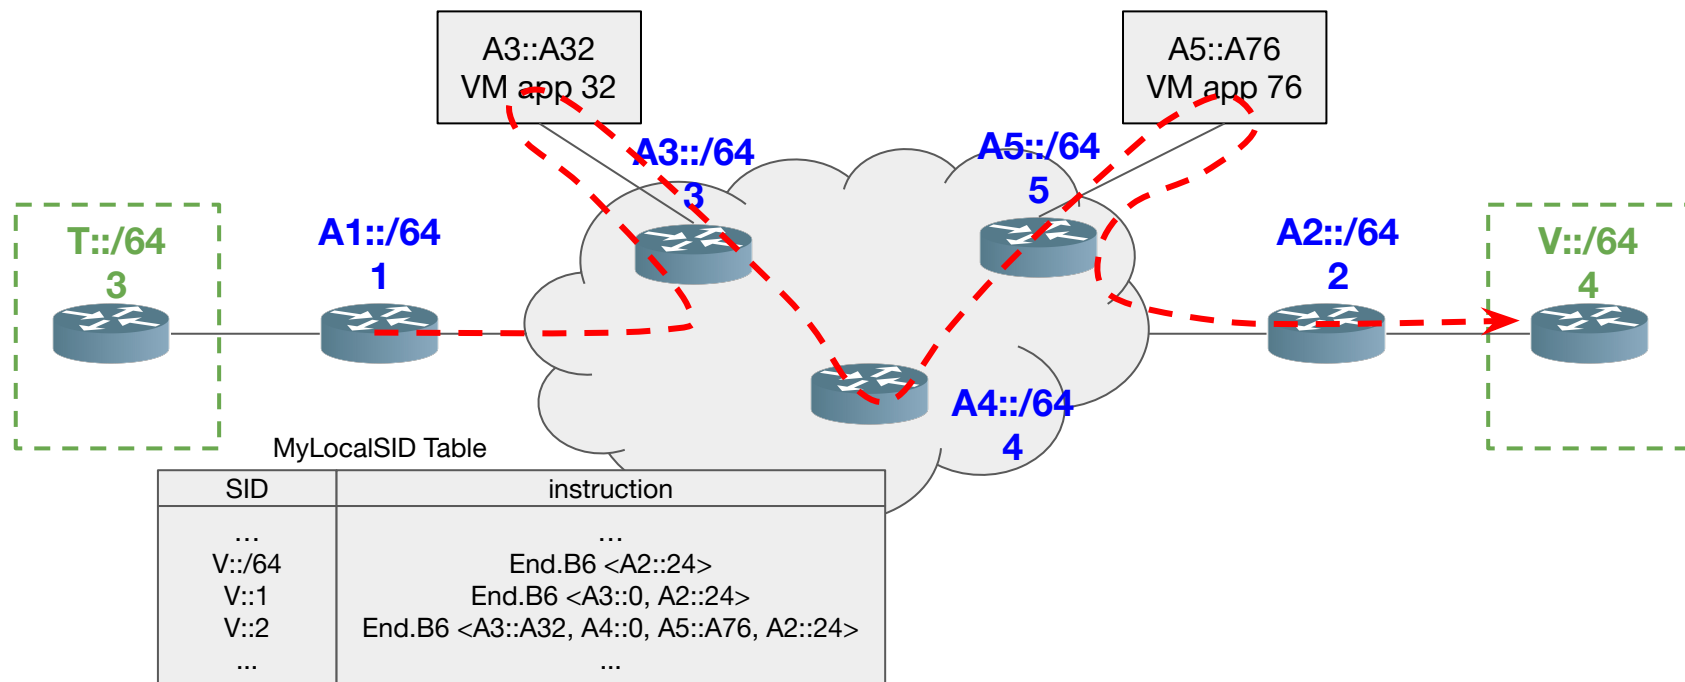




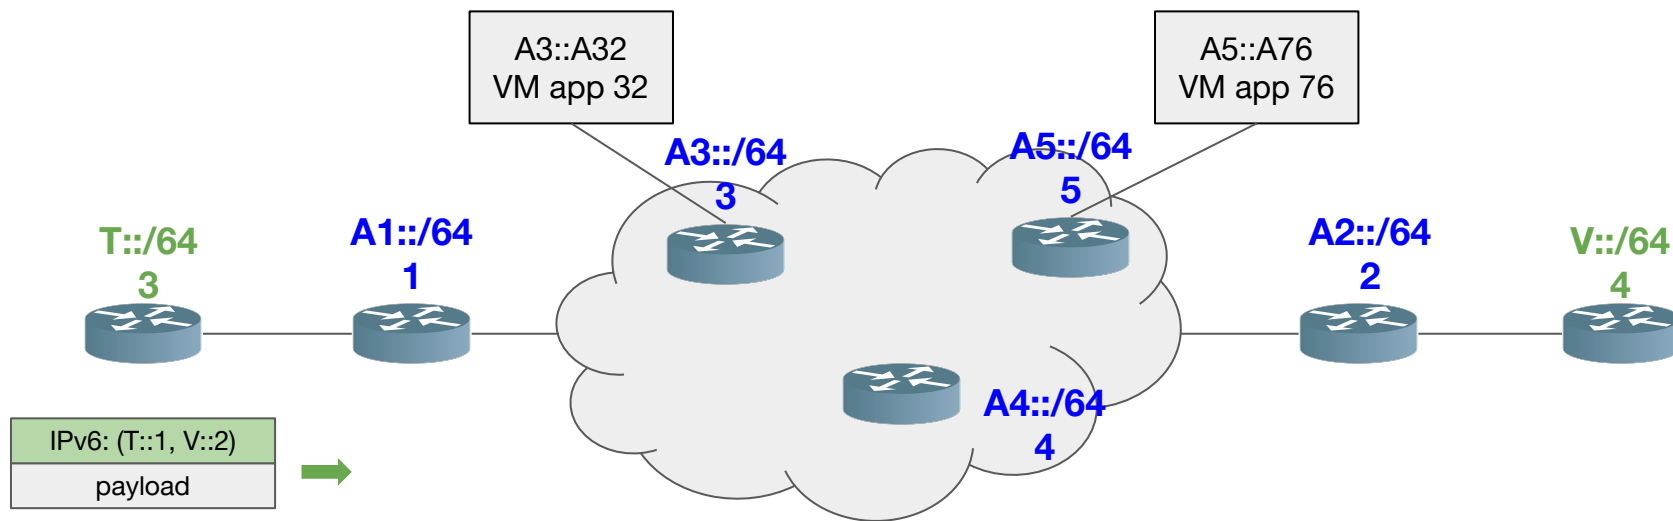
# Overlay with underlay SLA




# Integrated NFV



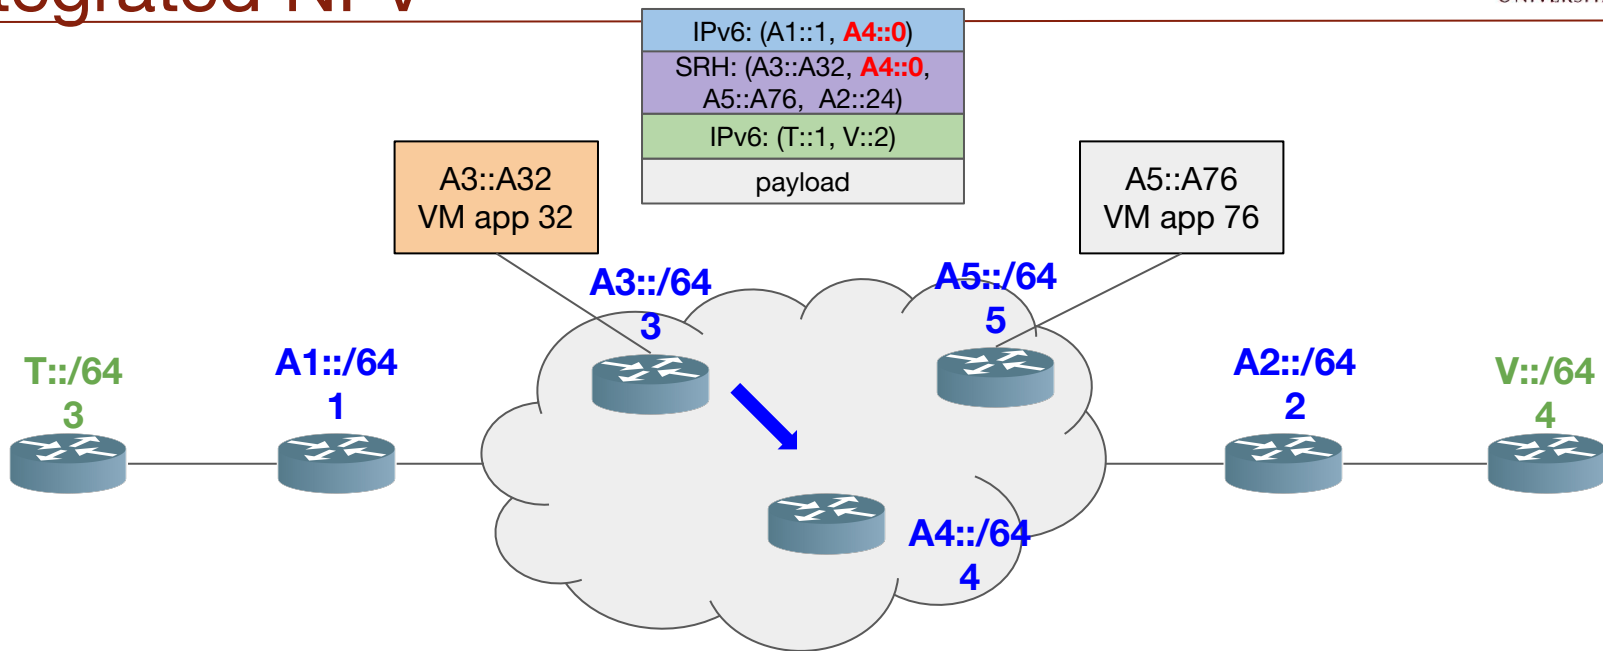
# Integrated NFV




# Integrated NFV



# Integrated NFV




# Integrated NFV




# Integrated NFV

IPv6: (A1::1, <b>A2::24</b> )
SRH: (A3::A32, A4::0, A5::A76, <b>A2::24</b> )
IPv6: (T::1, V::2)
payload



# Integrated NFV





# References

---

- James F. Kurose and Keith W. Ross. 2012. Computer Networking: A Top-Down Approach (6th Edition) (6th ed.). Pearson.
- <http://www.segment-routing.net/>
- C. Filsfils, N. K. Nainar, C. Pignataro, J. C. Cardona and P. Francois, "The Segment Routing Architecture," *2015 IEEE Global Communications Conference (GLOBECOM)*, San Diego, CA, 2015, pp. 1-6.
- Segment Routing Architecture, draft-ietf-spring-segment-routing-14
- Segment Routing Policy for Traffic Engineering, draft-filsfils-spring-segment-routing-policy-04.txt
- IPv6 Segment Routing Header (SRH), draft-ietf-6man-segment-routing-header-07
- SRv6 Network Programming, draft-filsfils-spring-srv6-network-programming-03